

Problem Statement:

Build a program or service that accepts:

1. An **audio file**
2. A **topic/question**
3. A **set of parameters**

The service should **analyze** the given audio file based on the specified parameters and return a **score for each parameter**.

Example Parameters:

1. **English Fluency** – Measures the speaker's fluency (e.g., speed, pauses, hesitations).
 2. **Vocabulary** – Evaluates the richness and diversity of words used.
 3. **Grammar** – Checks grammatical correctness.
 4. **Relevance to the Topic** – Assesses how closely the response aligns with the given topic.
-

Requirements & Guidelines:

- You **can use any model/service** (e.g., OpenAI Whisper, AWS Transcribe, Google Speech-to-Text, or custom NLP models).
 - The final output should be a **JSON response** with **parameter-wise scores**.
 - Provide a **README** with setup & execution instructions.
 - Submit a **demo video** (screen recording is fine).
 - The solution should be runnable on a **local machine** or a **server**.
-

Expected Output Example:

```
{  
  "fluency": 8.5,  
  "vocabulary": 7.2,  
  "grammar": 9.0,  
  "relevance": 6.8  
}
```

Submission Checklist:

- ✓ **Codebase** (GitHub repo or ZIP)
- ✓ **README** (setup, dependencies, execution steps)
- ✓ **Demo video** (showing end-to-end execution)

Bonus Points:

- Deploy the service as an **API** and share a live URL